



Churn Prediction & XG Boost Model Parameters

by Alex Dance

Purpose:

- Look at a Churn Model
- Explain the XG Boost parameters
- See how parameters can be changed

Agenda:

Part 1 - Research

- Explain Parameter Changing
- Tips I liked

Part 2 - Practice

- My dataset
- EDA
- How to run parameters

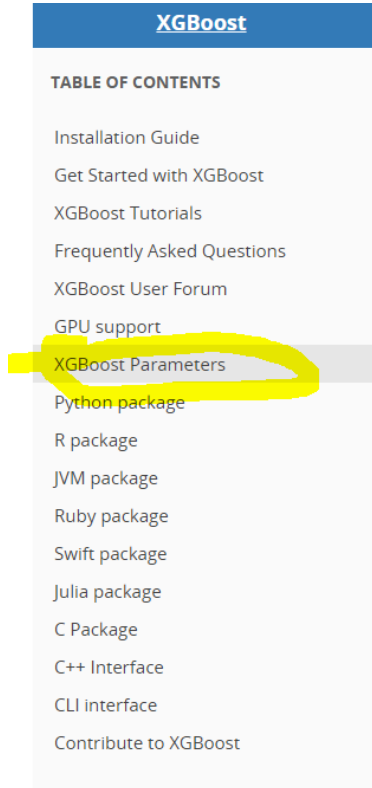
So your model does this



and NOT this



There are 3 main types of parameters



XGBoost
TABLE OF CONTENTS
Installation Guide
Get Started with XGBoost
XGBoost Tutorials
Frequently Asked Questions
XGBoost User Forum
GPU support
XGBoost Parameters
Python package
R package
JVM package
Ruby package
Swift package
Julia package
C Package
C++ Interface
CLI interface
Contribute to XGBoost

- **General parameters** relate to which booster we are using to do boosting, commonly tree or linear model
- **Booster parameters** depend on which booster you have chosen
- **Learning task parameters** decide on the learning scenario. For example, regression tasks may use different parameters with ranking tasks.
- **Command line parameters** relate to behavior of CLI version of XGBoost. (NOT relevant)

<https://xgboost.readthedocs.io/en/latest/parameter.html>

There are a few ways of setting up parameters

XGBoost can use either a list of pairs or a dictionary to set [parameters](#). For instance:

- Booster parameters

```
param = {'max_depth': 2, 'eta': 1, 'objective': 'binary:logistic'}  
param['nthread'] = 4  
param['eval_metric'] = 'auc'
```

- You can also specify multiple eval metrics:

```
param['eval_metric'] = ['auc', 'ams@0']  
  
# alternatively:  
# plst = param.items()  
# plst += [('eval_metric', 'ams@0')]
```

- Specify validations set to watch performance

```
evallist = [(dtest, 'eval'), (dtrain, 'train')]
```

https://xgboost.readthedocs.io/en/latest/python/python_intro.html#setting-parameters

Tip: a saved model can be exported and loaded

Training

Training a model requires a parameter list and data set.

```
num_round = 10  
bst = xgb.train(param, dtrain, num_round, evallist)
```

After training, the model can be saved.

```
bst.save_model('0001.model')
```

The model and its feature map can also be dumped to a text file.

```
# dump model  
bst.dump_model('dump.raw.txt')  
# dump model with feature map  
bst.dump_model('dump.raw.txt', 'featmap.txt')
```

A saved model can be loaded as follows:

```
bst = xgb.Booster({'nthread': 4}) # init model  
bst.load_model('model.bin') # load data
```

Methods including `update` and `boost` from `xgboost.Booster` are designed for internal usage only. The wrapper function `xgboost.train` does some pre-configuration including setting up caches and some other parameters.

Tip: Early stopping can still yield results

A model that has been trained or loaded can perform predictions on data sets.

```
# 7 entities, each contains 10 features
data = np.random.rand(7, 10)
dtest = xgb.DMatrix(data)
ypred = bst.predict(dtest)
```

If early stopping is enabled during training, you can get predictions from the best iteration with

```
bst.best_ntree_limit:
```

```
ypred = bst.predict(dtest, ntree_limit=bst.best_ntree_limit)
```

If early stopping occurs, the model will have three additional fields: `bst.best_score`, `bst.best_iteration` and `bst.best_ntree_limit`. Note that `xgboost.train()` will return a model from the last iteration, not the best one.

This works with both metrics to minimize (RMSE, log loss, etc.) and to maximize (MAP, NDCG, AUC). Note that if you specify more than one evaluation metric the last one in `param['eval_metric']` is used for early stopping.

Tip: Parameters running

1. Verbosity of printing messages. Valid values are

0 (silent)

1 (warning)

2 (info)

3 (debug)

nthread [default to maximum number of threads available if not set]

1. This is used for parallel processing and number of cores in the system should be entered

2. If you wish to run on all cores, value should not be entered and algorithm will detect automatically

Booster Parameters

From Vidhya

<https://www.analyticsvidhya.com/blog/2016/03/complete-guide-parameter-tuning-xgboost-with-codes-python/>

eta [default=0.3]

Analogous to learning rate in GBM

Makes the model more robust by shrinking the weights on each step

Typical final values to be used: 0.01-0.2

min_child_weight [default=1]

Defines the minimum sum of weights of all observations required in a child.

This is similar to **min_child_leaf** in GBM but not exactly. This refers to min "sum of weights" of observations while GBM has min "number of observations"

Used to control over-fitting. Higher values prevent a model from learning relations which might be highly specific to the particular sample selected for a tree.

Too high values can lead to under-fitting hence, it should be tuned using CV.

max_depth [default=6]

The maximum depth of a tree, same as GBM.

Used to control over-fitting as higher depth will allow model to learn relations very specific to a particular sample.

Should be tuned using CV.

Typical values: 3-10

max_leaf_nodes

The maximum number of terminal nodes or leaves in a tree.

Can be defined in place of max_depth. Since binary trees are created, a depth of 'n' would produce a maximum of 2^n leaves.

If this is defined, GBM will ignore max_depth.

gamma [default=0]

A node is split only when the resulting split gives a positive reduction in the loss function. Gamma specifies the minimum loss reduction required to make a split.

Makes the algorithm conservative. The values can vary depending on the loss function and should be tuned.

max_delta_step [default=0]

In maximum delta step we allow each tree's weight estimation to be. If the value is set to 0, it means there is no constraint. If it is set to a positive value, it can help making the update step more conservative.

Usually this parameter is not needed, but it might help in logistic regression when class is extremely imbalanced.

This is generally not used but you can explore further if you wish.

subsample [default=1]

Same as the subsample of GBM. Denotes the fraction of observations to be randomly samples for each tree.

Lower values make the algorithm more conservative and prevents overfitting but too small values might lead to under-fitting.

Typical values: 0.5-1

colsample_bytree [default=1]

Similar to max_features in GBM. Denotes the fraction of columns to be randomly samples for each tree.

Typical values: 0.5-1

colsample_bylevel [default=1]

Denotes the subsample ratio of columns for each split, in each level.

I don't use this often because subsample and colsample_bytree will do the job for you. but you can explore further if you feel so.

lambda [default=1]

L2 regularization term on weights (analogous to Ridge regression)

This used to handle the regularization part of XGBoost. Though many data scientists don't use it often, it should be explored to reduce overfitting.

alpha [default=0]

L1 regularization term on weight (analogous to Lasso regression)

Can be used in case of very high dimensionality so that the algorithm runs faster when implemented

scale_pos_weight [default=1]

A value greater than 0 should be used in case of high class imbalance as it helps in faster convergence.

Learning Task Parameters

These parameters are used to define the optimization objective the metric to be calculated at each step.

objective [default=reg:linear]

This defines the loss function to be minimized. Mostly used values are:

binary:logistic –logistic regression for binary classification, returns predicted probability (not class)

multi:softmax –multiclass classification using the softmax objective, returns predicted class (not probabilities)

you also need to set an additional **num_class** (number of classes) parameter defining the number of unique classes

multi:softprob –same as softmax, but returns predicted probability of each data point belonging to each class.

eval_metric [default according to objective]

The metric to be used for validation data.

The default values are rmse for regression and error for classification.

Typical values are:

rmse – root mean square error

mae – mean absolute error

logloss – negative log-likelihood

error – Binary classification error rate (0.5 threshold)

merror – Multiclass classification error rate

mlogloss – Multiclass logloss

auc: Area under the curve

seed [default=0]

The random number seed.

Can be used for generating reproducible results and also for parameter tuning.

General approach for parameter tuning

Choose a relatively **high learning rate**. Generally a learning rate of 0.1 works but somewhere between 0.05 to 0.3 should work for different problems.

Determine the **optimum number of trees for this learning rate**. XGBoost has a very useful function called as “cv” which performs cross-validation at each boosting iteration and thus returns the optimum number of trees required.

Tune tree-specific parameters (max_depth, min_child_weight, gamma, subsample, colsample_bytree) for decided learning rate and number of trees. Note that we can choose different parameters to define a tree and I'll take up an example here.

Tune **regularization parameters** (lambda, alpha) for xgboost which can help reduce model complexity and enhance performance.

Tip: There is a good guide to XG Boost on GitHub

Branch: master ▾


xgboost / demo / guide-python /

Create new file

Upload files


Find file

History

 trivialfis Implement Python data handler. (#5689) ...

✓ Latest commit 5af8161 23 days ago


..



README.md

[Breaking] Set output margin to True for custom objective. (#5564)


2 months ago



basic_walkthrough.py

Update Python demos with tests. (#5651)


last month



boost_from_prediction.py

Update Python demos with tests. (#5651)


last month



cross_validation.py

Update Python demos with tests. (#5651)


last month



custom_objective.py

Update Python demos with tests. (#5651)


last month



custom_rmsle.py

[Breaking] Set output margin to True for custom objective. (#5564)


2 months ago



custom_softmax.py

[Breaking] Set output margin to True for custom objective. (#5564)


2 months ago



evals_result.py

Update Python demos with tests. (#5651)


last month



external_memory.py

Update Python demos with tests. (#5651)

last month



gamma_regression.py

Update Python demos with tests. (#5651)

last month

<https://github.com/dmlc/xgboost/tree/master/demo/guide-python>

Now

my

code

Churn and Acquisition are regular tasks

<https://www.kaggle.com/blastchar/telco-customer-churn>



```
KaggleIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
```

Cleaning the data

```
In [8]: 1 mapping = {'Yes':1, 'No':0}
        2 df['Churn'] = df['Churn'].map(mapping)
        3
```

```
In [9]: 1 df['Churn'].value_counts()
```

```
Out[9]: 0    5174
        1    1869
        Name: Churn, dtype: int64
```

```
] 1 df['TotalCharges'] = pd.to_numeric(df['TotalCharges'],errors='coerce')
```

```
] 1 type(df['MonthlyCharges'])
```

```
1 df = pd.get_dummies(data = df, columns = ['gender', 'Partner','Dependents','PhoneService','MultipleLines','InternetService',
```

```
1 df.head()
```


customerID	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partner_No	Partner_Yes	...	PaperlessBilling_Yes
7590-VHVEG	0	1	29.85	29.85	0	1	0	0	1	...	1
5575-GNVDE	0	34	56.95	1889.5	0	0	1	1	0	...	0
3868-QPYBK	0	2	53.85	108.15	1	0	1	1	0	...	1

Looked at the data

MonthlyCharges
Real number ($\mathbb{R}_{\geq 0}$)

Distinct count	1585
Unique (%)	22.5%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%


Mean	64.76169246056
Minimum	18.25
Maximum	118.75
Zeros	0
Zeros (%)	0.0%
Memory size	55.1 KiB



tenure
Real number ($\mathbb{R}_{\geq 0}$)

Distinct count	73
Unique (%)	1.0%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%

Mean	32.37114865824
Minimum	0
Maximum	72
Zeros	11
Zeros (%)	0.2%
Memory size	55.1 KiB




Toggle details

TotalCharges
Real number ($\mathbb{R}_{\geq 0}$)

Distinct count	6531
Unique (%)	92.7%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%

Mean	2283.30044084
Minimum	18.8
Maximum	8684.8
Zeros	0
Zeros (%)	0.0%
Memory size	55.1 KiB



Looking at Dependent Variables

```
1 df.groupby(["Churn"])["tenure"].mean()
```

```
Churn
0    37.569965
1    17.979133
Name: tenure, dtype: float64
```

```
1 df.groupby(["Churn"])["Partner_No"].mean()
```

```
Churn
0    0.471782
1    0.642055
Name: Partner_No, dtype: float64
```

```
1 df.groupby(["Churn"])["TotalCharges"].mean()
```

```
Churn
0    2555.344141
1    1531.796094
Name: TotalCharges, dtype: float64
```



More likely to churn if

Tenure was LOWER

Spend was Lower

Ran Oversampling

```
train_b2, test_b2 = train_test_split(df, test_size=0.20, random_state=1)
```

```
1 target_count = train_b2.Churn.value_counts()
```

```
1 print(target_count)
```

```
0    4113  
1    1521  
Name: Churn, dtype: int64
```

```
1 count_class_1 = target_count[1]  
2 count_class_0 = target_count[0]
```

```
In [71]: 1 df_class_0 = train_b2[train_b2['Churn'] == 0]  
2 df_class_1 = train_b2[train_b2['Churn'] == 1]
```

```
In [72]: 1 # Random over-sampling  
2 df_class_1_over = df_class_1.sample(count_class_0, replace=True)  
3 df_train_over = pd.concat([df_class_0, df_class_1_over], axis=0)
```

```
In [74]: 1 df_train_over['Churn'].value_counts()
```

```
Out[74]: 1    4113  
0    4113  
Name: Churn, dtype: int64
```

```
In [75]: 1 # Get all variables
```

```
In [76]: 1 y_test_bs = test_b2 ['Churn']  
2 X_test_bs = test_b2.drop(["customerID", "Churn"], axis = 1)  
3  
4 y_train_bs = df_train_over[['Churn']]  
5 X_train_bs = df_train_over.drop(["customerID", "Churn"], axis = 1)  
6
```

4 Basic XG Boost Models

Default Limited Features

XG_Results Results

```
XG_Results AUC Test 79.69%  
[[965 96]  
[199 149]]  
XG_Results accuracy_score_train 77.65%  
XG_Results accuracy_score_test 79.06%
```

All – except - top 3 features

```
XG_Results AUC Test 79.55%  
[[878 183]  
[153 195]]  
XG_Results accuracy_score_train 88.62%  
XG_Results accuracy_score_test 76.15%
```

```
X_b = df.drop(["customerID","Churn"], axis = 1)
```

All With oversampling

```
xG train score 94.98%  
XG test score 76.30%
```

All - Standard

XG_Results Results

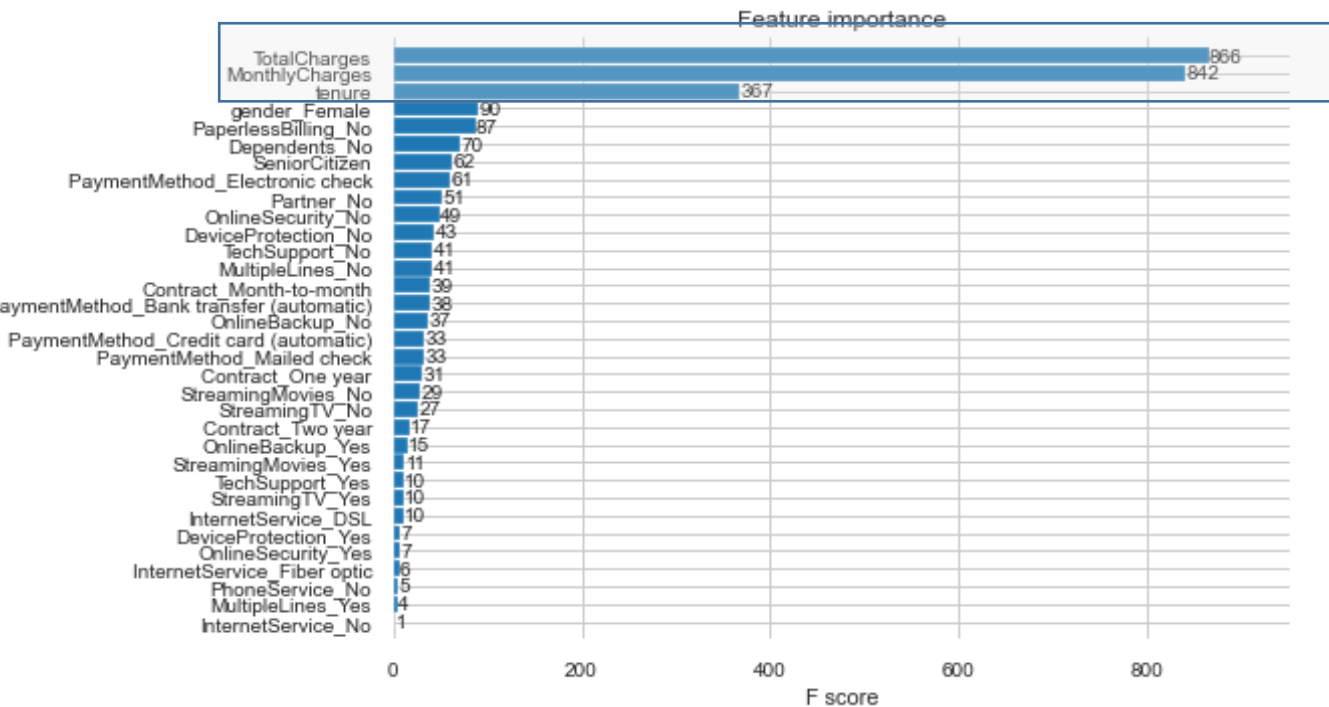
```
XG_Results AUC Test 83.74%  
[[926 135]  
[147 201]]  
XG_Results accuracy_score_train 93.81%  
XG_Results accuracy_score_test 79.99%
```

Plus Logistic regression - limited

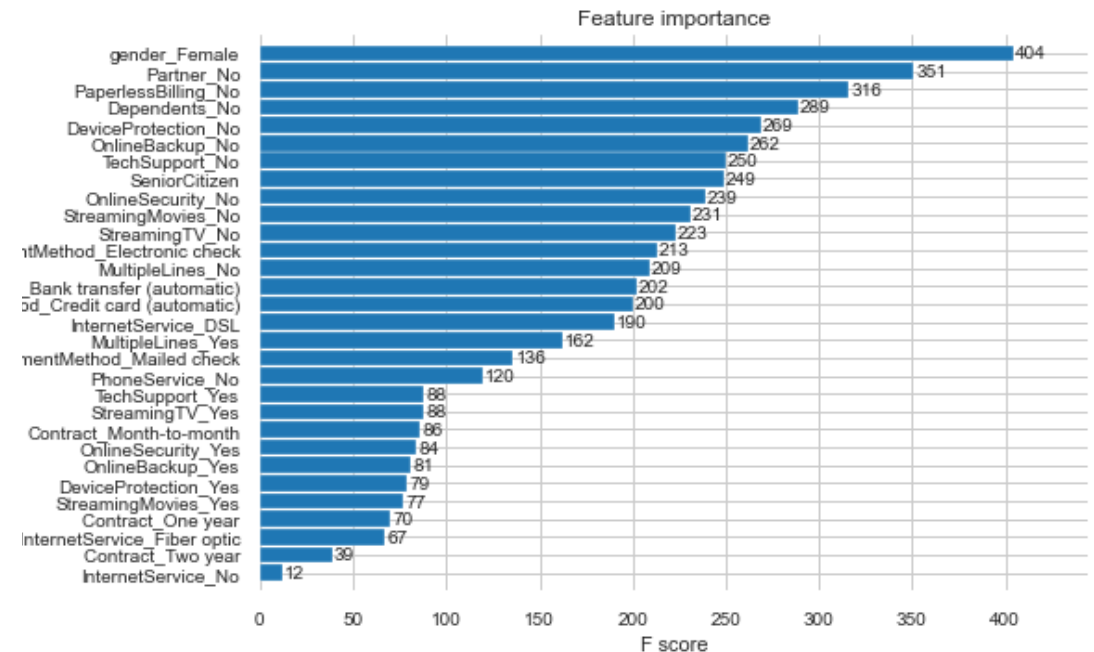
```
train score 77%  
test score 79%
```

Looked at the feature importance

Original



Updated



Could be Negatively or Positively Correlated

Grid Search CV

```
1 from sklearn.model_selection import GridSearchCV
```

```
1 xgb_modelcv = XGBClassifier(params = params)
```

```
1 test_params = {  
2     'max_depth': [4, 8, 12]  
3 }  
4
```

```
1 model = GridSearchCV(estimator = xgb_modelcv, param_grid = test_params)
```

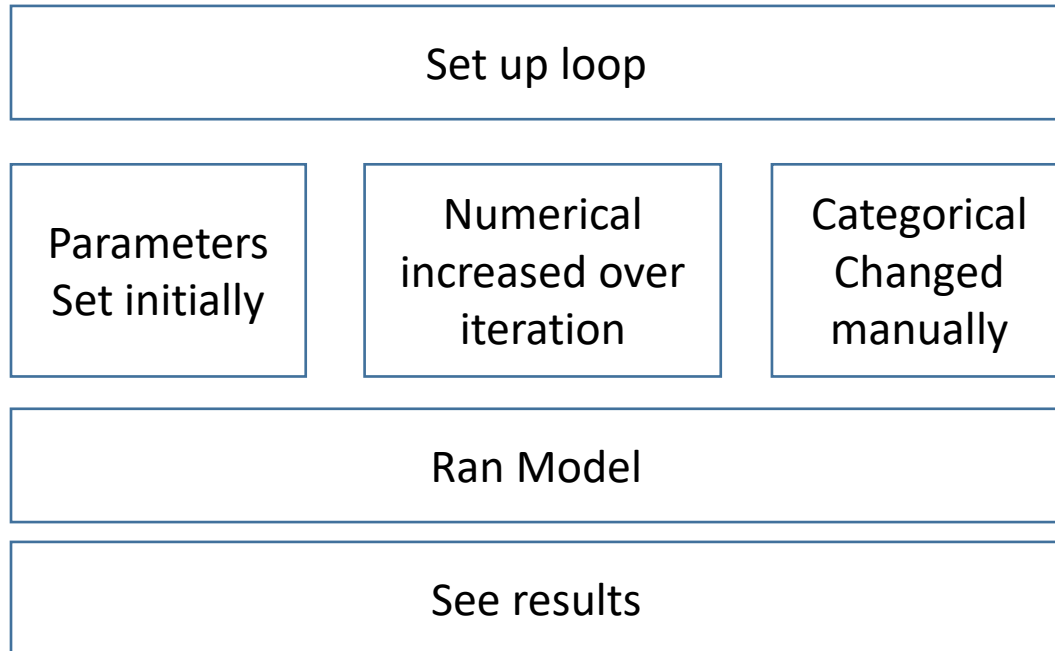
```
1 model.fit(X_train_b, y_train_b)  
2
```

In [117]:

```
1 test_params = {  
2     'max_depth': [2, 4, 8, 12]  
3 }  
4
```

```
: 1 print (model.best_params_)  
{'max_depth': 2}
```

Iterated over Function to check everything



7 Loops

12 Variables

Over 50 different tests

Was able to change multiple parameters

Can comment out all but 1 variable

Start

```
LearningRate = 0.1
MaxDepth = 1
Alpha = 1
Grow_Policy = 'lossguide'
ColSampleByTree = 0.1
MaximumDepth = 1 #'alpha'
Colsample_bytree = 0.3
Colsample_bylevel = 0.3
Colsample_bynode = 0.3
TreeMethod = 'auto'
Process_type = 'default'
```

Changes

```
# Every time we go through change values
LearningRate += 0.1
MaxDepth +=1
ColSampleByTree +=0.1
Alpha +=1
Colsample_bytree = Colsample_bytree +0.1
Colsample_bylevel = Colsample_bylevel +0.1
Colsample_bynode = Colsample_bynode +0.1
if i == 1:
    MaximumDepth = 1
    TreeMethod = 'exact'

if i > 1:
    Grow_Policy = 'depthwise'
    MaximumDepth += 1
    TreeMethod = 'approx'
    Updater = 'prune'
```

```
# Every time we go through change values
LearningRate += 0.1
#MaxDepth +=1
#ColSampleByTree +=0.1
#Alpha +=1
#Colsample_bytree = Colsample_bytree +0.1
#Colsample_bylevel = Colsample_bylevel +0.1
#Colsample_bynode = Colsample_bynode +0.1
if i == 1:
    irrelevant = 1
    # MaximumDepth = 1
    # TreeMethod = 'exact'


if i > 1:
    irrelevant = 2
    # Grow_Policy = 'depthwise'
    # MaximumDepth += 1
    # TreeMethod = 'approx'
    # Updater = 'prune'
```

Multiple Parameters changed

```
def find_all(y_test_f,X_test_f,model_f,X_train_f, y_train_f):  
    model_f.fit(X_train_f, y_train_f)  
    preds = model_f.predict_proba(X_test_f)[:,-1]  
    fpr, tpr, thresholds = metrics.roc_curve(y_test_f, preds)  
    roc_auc = metrics.auc(fpr, tpr)  
    y_pred_f = model_f.predict(X_test_f)  
    cf = confusion_matrix(y_test_f, y_pred_f)  
    accuracy_score_train = model_f.score(X_train_f, y_train_f)  
    accuracy_score_test = model_f.score(X_test_f, y_test_f)  
    return{'auc': roc_auc, 'cfm':cf , 'accuracy_score_train':accuracy_score_train , 'accuracy_score_test':accuracy_score_test}
```

For each iteration printed

```
## Below is Printing
print(color.BOLD + " ", i ,") XG_Results")
print ('\033[0m')
print(" ")
print( "learning_rate", LearningRate)
print("process_type ", Process_type)
print( "Updater ", Updater)
print("max_depth", MaxDepth)
print( "grow_policy", Grow_Policy)
print("tree_method", TreeMethod)
print( "colsample_bytree", Colsample_bytree, "colsample_bylevel", Colsample_bylevel , "colsample_bynode", Colsample_bynode)
print(" ")
print("XG_Results AUC Test  %.2f%%" % (XG_Results2['auc']* 100.0))
print(XG_Results2['cfm'])
print("XG_Results accuracy_score_train  %.2f%%" % (XG_Results2 ['accuracy_score_train'] * 100.0))
print("XG_Results accuracy_score_test  %.2f%%" % (XG_Results2 ['accuracy_score_test']* 100.0))
print(" ")
```



4) XG_Results

```
learning_rate 0.4
process_type default
max_depth 4
grow_policy depthwise
tree_method approx
colsample_bytree 0.6 colsample_bylevel 0.6 colsample_bynode 0.6
```

```
XG_Results AUC Test  81.88%
[[913 148]
 [157 191]]
XG_Results accuracy_score_train  81.66%
XG_Results accuracy_score_test   78.35%
```


Ran some AB tests to see what changed

NO CHANGE

grow_policy= 'lossguide' -> depthwise
tree_method= 'auto' - > exact -> approx
colsample_bytree= 0.3 -> 1
colsample_bynode= 0.3 - >0.6

RESULTS CHANGE

colsample_bytree= 1 -> 0.5 AUC down 0.1%
learning_rate= 0.2 -> 0.4 AUC up 0.5%
colsample_bylevel= 0.3 - >0.6 AUC down 0.1%

Eg (colsample_bytree= 0.5, learning_rate= 0.1, max_depth= 2, grow_policy= 'depthwise', colsample_bylevel= 0.3, colsample_bynode= 0.6, tree_method= 'approx', process_type= 'default', verbosity = 1)

Can see impact of all changes OR 1 at a time

1) XG_Rsults

```
learning_rate 0.1
process_type default
max_depth 1
grow_policy lossguide
tree_method auto
colsample_bytree 0.3 colsample_bylevel 0.3 colsample_bynode 0.3

XG_Rsults AUC Test 83.80%
[[1009 52]
 [232 116]]
XG_Rsults accuracy_score_train 76.89%
XG_Rsults accuracy_score_test 79.84%
```

2) XG_Rsults

```
learning_rate 0.2
process_type default
max_depth 2
grow_policy lossguide
tree_method exact
colsample_bytree 0.4 colsample_bylevel 0.4 colsample_bynode 0.4

XG_Rsults AUC Test 84.44%
[[941 120]
 [159 189]]
XG_Rsults accuracy_score_train 79.02%
XG_Rsults accuracy_score_test 80.20%
```

5) XG_Rsults

```
learning_rate 0.5
process_type default
max_depth 5
grow_policy depthwise
tree_method approx
colsample_bytree 0.7 colsample_bylevel 0.7 colsample_bynode 0.7

XG_Rsults AUC Test 79.18%
[[884 177]
 [153 195]]
XG_Rsults accuracy_score_train 86.39%
XG_Rsults accuracy_score_test 76.58%
```

6) XG_Rsults

```
learning_rate 0.6
process_type default
max_depth 6
grow_policy depthwise
tree_method approx
colsample_bytree 0.7999999999999999 colsample_bylevel 0.7999999999999999 colsample_bynode 0.7999999999999999

XG_Rsults AUC Test 77.19%
[[862 199]
 [152 196]]
XG_Rsults accuracy_score_train 90.61%
XG_Rsults accuracy_score_test 75.09%
```

Changing 1 Variable

2) XG_Rsults

```
learning_rate 0.2
process_type default
max_depth 1
grow_policy lossguide
tree_method auto
colsample_bytree 0.3 colsample_bylevel 0.3 colsample_bynode 0.3
```

```
XG_Rsults AUC Test 84.41%
[[956 105]
 [170 178]]
XG_Rsults accuracy_score_train 78.04%
XG_Rsults accuracy_score_test 80.48%
```

3) XG_Rsults

```
learning_rate 0.30000000000000004
process_type default
max_depth 1
grow_policy lossguide
tree_method auto
colsample_bytree 0.3 colsample_bylevel 0.3 colsample_bynode 0.3
```

```
XG_Rsults AUC Test 84.62%
[[939 122]
 [162 186]]
XG_Rsults accuracy_score_train 78.24%
XG_Rsults accuracy_score_test 79.84%
```

4) XG_Rsults

```
learning_rate 0.4
process_type default
max_depth 1
grow_policy lossguide
tree_method auto
colsample_bytree 0.3 colsample_bylevel 0.3 colsample_bynode 0.3
```

```
XG_Rsults AUC Test 84.68%
[[931 130]
 [158 190]]
XG_Rsults accuracy_score_train 78.31%
XG_Rsults accuracy_score_test 79.56%
```

Best

5) XG_Rsults

```
learning_rate 0.5
process_type default
max_depth 1
grow_policy lossguide
tree_method auto
colsample_bytree 0.3 colsample_bylevel 0.3 colsample_bynode 0.3
```

```
XG_Rsults AUC Test 84.65%
[[930 131]
 [158 190]]
XG_Rsults accuracy_score_train 78.56%
XG_Rsults accuracy_score_test 79.49%
```

Starts getting worse

Thanks

Alex Dance



Background

- Maths / statistics degree
- Background in big data, strategy, analytics
- Worked at Optus, Salmat, Reuters, Pathfinder Solutions

Copy of This Presentation and code

<https://github.com/alexdance2468/>

Plus other data science projects completed

Contact Details

www.linkedin.com/in/alex-dance/